

Notebook: Comparing Segmentation models: Segment anything model (SAM) vs UNET on image segmentation

In this notebook, we'll reproduce the SAM project by MetaAI, we will look into the model architectures of Segment anything model and UNET which has been mostly used image segmentation, we will later fine train SAM and UNET on a dataset of medical images.

Team member 1: Anupam Tiwari (ast9885)

Team member 2: Ankur Aggarwal (aa10336)



Set-up environment

We first install Transformers, Datasets and other nessasry libraries that our project is dependent on

In [1]:

```
! pip install opencv-python pycocotools matplotlib onnxruntime onnx pycocotools supervision &> /dev/null
! pip install -q git+https://github.com/huggingface/transformers.git
! pip install -q datasets
! pip install git+https://github.com/facebookresearch/segment-anything.git &> /dev/null
! wget https://dl.fbaipublicfiles.com/segment_anything/sam_vit_b_01ec64.pth &> /dev/null
```

```
Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
_____ 224.5/224.5 kB 5.8 MB/s eta 0:00:00
_____ 7.8/7.8 MB 59.3 MB/s eta 0:00:00
Building wheel for transformers (pyproject.toml) ... done
_____ 474.6/474.6 kB 10.3 MB/s eta 0:00:00
_____ 110.5/110.5 kB 14.2 MB/s eta 0:00:00
_____ 212.5/212.5 kB 24.2 MB/s eta 0:00:00
_____ 134.3/134.3 kB 12.5 MB/s eta 0:00:00
_____ 1.0/1.0 MB 32.5 MB/s eta 0:00:00
_____ 114.5/114.5 kB 13.1 MB/s eta 0:00:00
_____ 268.8/268.8 kB 26.9 MB/s eta 0:00:00
_____ 149.6/149.6 kB 17.1 MB/s eta 0:00:00
```

importing dependencies

In [2]:

```
from pycocotools.coco import COCO
import os
from PIL import Image
import numpy as np
from matplotlib import pyplot as plt
import cv2
import json
from torch.utils.data import Dataset
import torch
import torchvision
from torch.utils.data import DataLoader
import albumentations as A
```

```
from albumentations.pytorch import ToTensorV2
from tqdm import tqdm
import torch.nn as nn
import torch.optim as optim
%matplotlib inline
```

Load dataset

Here we load a small dataset of 130 (image, ground truth mask) pairs. for SAM and UNET

In [3]:

```
from datasets import load_dataset

dataset = load_dataset("nielsr/breast-cancer", split="train")
```

Downloading and preparing dataset None/None to /root/.cache/huggingface/datasets/nielsr__parquet/nielsr--breast-cancer-c16ee7932c43ffa3/0.0.0/2a3b91fbd88a2c90d1dbbb32b460cf621d31bd5b05b934492fdef7d8d6f236ec...

Dataset parquet downloaded and prepared to /root/.cache/huggingface/datasets/nielsr__parquet/nielsr--breast-cancer-c16ee7932c43ffa3/0.0.0/2a3b91fbd88a2c90d1dbbb32b460cf621d31bd5b05b934492fdef7d8d6f236ec. Subsequent calls will reuse this data.

Lets look into the dataset: it has image and label pair, label is the mask

In [4]:

```
dataset
```

Out[4]:


```
Dataset({
  features: ['image', 'label'],
  num_rows: 130
})
```

example image:

In [5]:

```
example = dataset[0]
image = example["image"]
image
```

Out[5]:



IMAGE+MASK

In [6]:

```
import matplotlib.pyplot as plt
import numpy as np

def show_mask(mask, ax, random_color=False):
    if random_color:
        color = np.concatenate([np.random.random(3), np.array([0.6])], axis=0)
    else:
        color = np.array([30/255, 144/255, 255/255, 0.6])
    h, w = mask.shape[-2:]
    mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
    ax.imshow(mask_image)


fig, axes = plt.subplots()

axes.imshow(np.array(image))
ground_truth_seg = np.array(example["label"])
show_mask(ground_truth_seg, axes)
axes.title.set_text(f"Ground truth mask")
axes.axis("off")
```

Out[6]:

(-0.5, 255.5, 255.5, -0.5)

Ground truth mask




Mask:


In [10]:

```
image = dataset[0]["image"]
mask = np.array(dataset[0]["label"])
h, w = mask.shape[-2:]
color = np.array([30/255, 144/255, 255/255, 0.6])
mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
plt.imshow(mask_image)
```

Out[10]:

<matplotlib.image.AxesImage at 0x7f97849152d0>





Now we will make two directories for Mask and Image, this is need for the UNET

In [11]:

```
!mkdir mask
!mkdir img
```

```
mkdir: cannot create directory 'mask': File exists
mkdir: cannot create directory 'img': File exists
```

In [12]:

```
for i in range(130):
    image = np.array(dataset[i]["image"])
    mask = np.array(dataset[i]["label"])
    h, w = mask.shape[-2:]
    color = np.array([30/255, 144/255, 255/255, 0.6])
    mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
    file_name = str(i)
    cv2.imwrite('mask/'+file_name+".png", mask_image)
    cv2.imwrite('img/'+file_name+".png", image)
```

Create PyTorch dataset

Below we define a regular PyTorch dataset, which gives us examples of the data prepared in the format for the model. Each example consists of:

pixel values (which is the image prepared for the model) a prompt in the form of a bounding box a ground truth segmentation mask. The function below defines how to get a bounding box prompt based on the ground truth segmentation. This was taken from [here](#).

Note that SAM is always trained using certain "prompts", which you could be bounding boxes, points, text, or rudimentary masks. The model is then trained to output the appropriate mask given the image + prompt.

In [13]:

```
def get_bounding_box(ground_truth_map):
    # get bounding box from mask
    y_indices, x_indices = np.where(ground_truth_map > 0)
    x_min, x_max = np.min(x_indices), np.max(x_indices)
    y_min, y_max = np.min(y_indices), np.max(y_indices)
    # add perturbation to bounding box coordinates
    H, W = ground_truth_map.shape
    x_min = max(0, x_min - np.random.randint(0, 20))
    x_max = min(W, x_max + np.random.randint(0, 20))
    y_min = max(0, y_min - np.random.randint(0, 20))
    y_max = min(H, y_max + np.random.randint(0, 20))
    bbox = np.array([x_min, y_min, x_max, y_max])
```

```
return bbox
```

In [15]:

```
bbox_coords = {}  
for x in range(130):  
    bbox_coords[str(x)] = get_bounding_box(np.array(dataset[x]['label']))
```

In [16]:

```
ground_truth_masks = {}  
for k in bbox_coords.keys():  
    gt_grayscale = np.array(dataset[int(k)]['label'])  
    ground_truth_masks[k] = (gt_grayscale == 1)  
ground_truth_masks['1']
```

Out[16]:

```
array([[False, False, False, ..., False, False, False],  
       [False, False, False, ..., False, False, False],  
       [False, False, False, ..., False, False, False],  
       ...,  
       [False, False, False, ..., False, False, False],  
       [False, False, False, ..., False, False, False],  
       [False, False, False, ..., False, False, False]])
```

In [17]:

```
from torch.utils.data import Dataset  
  
class SAMDataset(Dataset):  
    def __init__(self, dataset, processor):  
        self.dataset = dataset  
        self.processor = processor  
  
    def __len__(self):  
        return len(self.dataset)  
  
    def __getitem__(self, idx):  
        item = self.dataset[idx]  
        image = item["image"]  
        ground_truth_mask = np.array(item["label"])  
  
        # get bounding box prompt  
        prompt = get_bounding_box(ground_truth_mask)  
  
        # prepare image and prompt for the model  
        inputs = self.processor(image, input_boxes=[[prompt]]), return_tensors="pt")  
  
        # remove batch dimension which the processor adds by default  
        inputs = {k:v.squeeze(0) for k,v in inputs.items()}  
  
        # add ground truth segmentation  
        inputs["ground_truth_mask"] = ground_truth_mask  
  
        return inputs
```

In [18]:

```
from transformers import SamProcessor  
  
processor = SamProcessor.from_pretrained("facebook/sam-vit-base")
```

In [19]:

```
train_dataset = SAMDataset(dataset=dataset, processor=processor)
```

In [20]:

```
example = train_dataset[0]
```

```
for k,v in example.items():
    print(k,v.shape)
```

```
pixel_values torch.Size([3, 1024, 1024])
original_sizes torch.Size([2])
reshaped_input_sizes torch.Size([2])
input_boxes torch.Size([1, 4])
ground_truth_mask (256, 256)
```

Take a look at the images, the bounding box prompts and the ground truth segmentation masks

In [21]:

```
# Helper functions provided in https://github.com/facebookresearch/segment-anything/blob/
9e8f1309c94f1128a6e5c047a10fdcb02fc8d651/notebooks/predictor_example.ipynb
def show_mask(mask, ax, random_color=False):
    if random_color:
        color = np.concatenate([np.random.random(3), np.array([0.6])], axis=0)
    else:
        color = np.array([30/255, 144/255, 255/255, 0.6])
    h, w = mask.shape[-2:]
    mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
    ax.imshow(mask_image)


def show_box(box, ax):
    print(box)
    x0, y0 = box[0], box[1]
    w, h = box[2] - box[0], box[3] - box[1]
    ax.add_patch(plt.Rectangle((x0, y0), w, h, edgecolor='green', facecolor=(0,0,0,0), 1
w=2))
```


In [22]:

```
name = '1'
image = cv2.imread(f'img/{name}.png')

plt.figure(figsize=(10,10))
plt.imshow(image)
show_box(bbox_coords[name], plt.gca())
show_mask(ground_truth_masks[name], plt.gca())
plt.axis('off')
plt.show()
```

```
[ 65  39 160 137]
```





Prepare Fine Tuning

Create PyTorch DataLoader


Next we define a PyTorch Dataloader, which allows us to get batches from the dataset.

In [23]:

```
from torch.utils.data import DataLoader

train_dataloader = DataLoader(train_dataset, batch_size=2, shuffle=True)
```

Load Model



In [24]:

```
model_type = 'vit_b'
checkpoint = 'sam_vit_b_01ec64.pth'
device = 'cuda:0'
```

In [25]:

```
from transformers import SamModel

sam_model = SamModel.from_pretrained("facebook/sam-vit-base")

# make sure we only compute gradients for mask decoder
for name, param in sam_model.named_parameters():
    if name.startswith("vision_encoder") or name.startswith("prompt_encoder"):
        param.requires_grad_(False)
```

In [26]:

```
# Set up the optimizer, hyperparameter tuning will improve performance here
lr = 1e-4
wd = 0
optimizer = torch.optim.Adam(sam_model.mask_decoder.parameters(), lr=lr, weight_decay=wd)

loss_fn = torch.nn.MSELoss()
# loss_fn = torch.nn.BCELoss()
keys = list(bbox_coords.keys())
```


Run fine tuning

This is the main training loop.

Improvements to be made include batching and moving the computation of the image and prompt embeddings outside the loop since we are not tuning these parts of the model, this will speed up training as we should not recompute the embeddings during each epoch.

In a production implementation a better choice of optimiser/loss function will certainly help.

using monai for custom loss function, called dice loss



In [28]:

```
!pip install monai
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
Requirement already satisfied: monai in /usr/local/lib/python3.10/dist-packages (1.1.0)
Requirement already satisfied: torch>=1.8 in /usr/local/lib/python3.10/dist-packages (from monai) (2.0.0+cu118)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from monai) (1.22.4)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.8->monai) (3.12.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from torch>=1.8->monai) (4.5.0)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.8->monai) (1.11.1)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.8->monai) (3.1)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8->monai) (3.1.2)
Requirement already satisfied: triton==2.0.0 in /usr/local/lib/python3.10/dist-packages (


```
from torch>=1.8->monai) (2.0.0)
Requirement already satisfied: cmake in /usr/local/lib/python3.10/dist-packages (from triton==2.0.0->torch>=1.8->monai) (3.25.2)
Requirement already satisfied: lit in /usr/local/lib/python3.10/dist-packages (from triton==2.0.0->torch>=1.8->monai) (16.0.3)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.8->monai) (2.1.2)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy->torch>=1.8->monai) (1.3.0)
```

In [29]:

```
from torch.optim import Adam
import monai

# Note: Hyperparameter tuning could improve performance here
optimizer = Adam(sam_model.mask_decoder.parameters(), lr=1e-5, weight_decay=0)

seg_loss = monai.losses.DiceCELoss(sigmoid=True, squared_pred=True, reduction='mean')
```

In [38]:

```
from tqdm import tqdm
from statistics import mean
import torch
from torch.nn.functional import threshold, normalize

num_epochs = 10

device = "cuda" if torch.cuda.is_available() else "cpu"
sam_model.to(device)

sam_model.train()
losses = []
for epoch in range(num_epochs):
    epoch_losses = []
    for batch in tqdm(train_dataloader):
        # forward pass
        outputs = sam_model(pixel_values=batch["pixel_values"].to(device),
                             input_boxes=batch["input_boxes"].to(device),
                             multimask_output=False)

        # compute loss
        predicted_masks = outputs.pred_masks.squeeze(1)
        ground_truth_masks = batch["ground_truth_mask"].float().to(device)
        loss = seg_loss(predicted_masks, ground_truth_masks.unsqueeze(1))

        # backward pass (compute gradients of parameters w.r.t. loss)
        optimizer.zero_grad()
        loss.backward()

        # optimize
        optimizer.step()
        epoch_losses.append(loss.item())

    print(f'EPOCH: {epoch}')
    print(f'Mean loss: {mean(epoch_losses)}')
    losses.append(mean(epoch_losses))
```

```
100%|██████████| 65/65 [01:03<00:00, 1.02it/s]
```

```
EPOCH: 0
Mean loss: 0.09859017454660855
```

```
100%|██████████| 65/65 [01:02<00:00, 1.04it/s]
```

```
EPOCH: 1
Mean loss: 0.09546073858554546
```

```
100%|██████████| 65/65 [01:02<00:00, 1.03it/s]
```

```
EPOCH: 2
Mean loss: 0.09198274795825664
```

Mean loss: 0.07983663311371436

100%|██████████| 65/65 [01:02<00:00, 1.04it/s]

EPOCH: 3

Mean loss: 0.08765292396912208

100%|██████████| 65/65 [01:02<00:00, 1.03it/s]

EPOCH: 4

Mean loss: 0.08331799965638381

100%|██████████| 65/65 [01:02<00:00, 1.04it/s]

EPOCH: 5

Mean loss: 0.07983663311371436

100%|██████████| 65/65 [01:03<00:00, 1.03it/s]

EPOCH: 6

Mean loss: 0.07851512294549208

100%|██████████| 65/65 [01:02<00:00, 1.03it/s]

EPOCH: 7

Mean loss: 0.08010301269017733

100%|██████████| 65/65 [01:02<00:00, 1.04it/s]

EPOCH: 8

Mean loss: 0.07710946844174311

100%|██████████| 65/65 [01:02<00:00, 1.03it/s]

EPOCH: 9

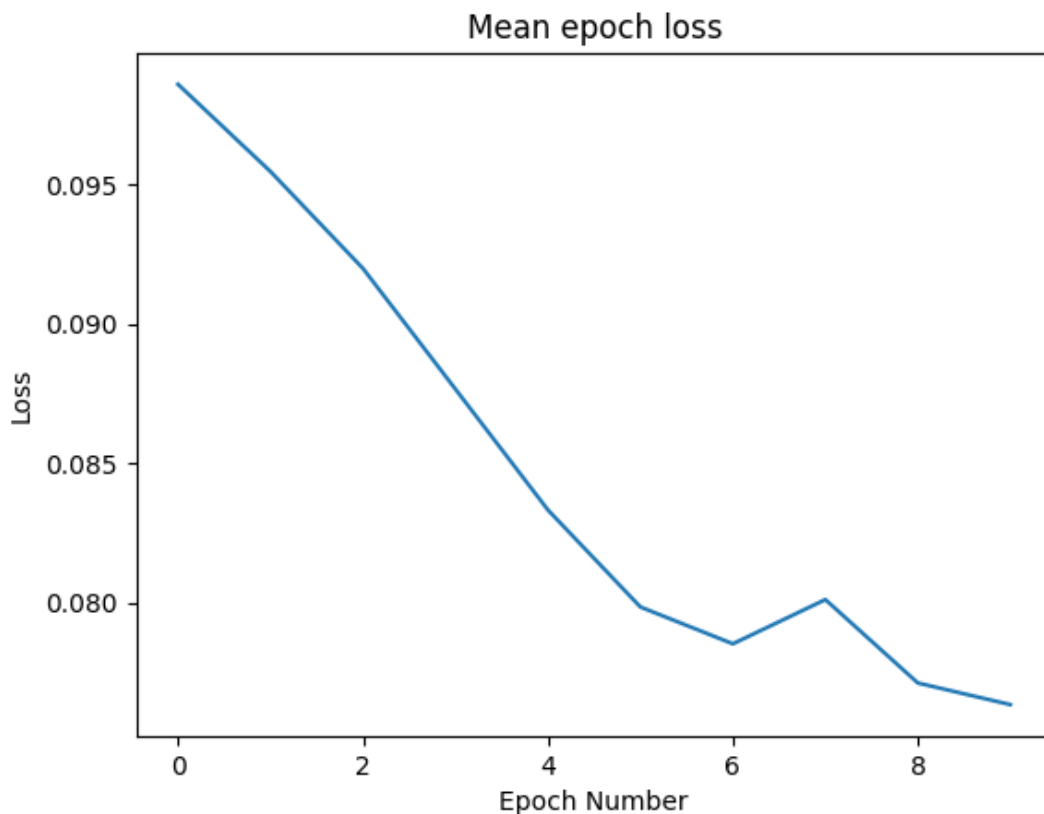
Mean loss: 0.07633395791053772

In [39]:

```
mean_losses = losses

plt.plot(list(range(len(mean_losses))), mean_losses)
plt.title('Mean epoch loss')
plt.xlabel('Epoch Number')
plt.ylabel('Loss')

plt.show()
```



We can compare our tuned model to the original model


In [40]:

```
import numpy as np
from PIL import Image

# let's take a random training example
idx = 10

# load image
image = dataset[idx]["image"]
image
```

Out[40]:



In [41]:

```
# get box prompt based on ground truth segmentation map
ground_truth_mask = np.array(dataset[idx]["label"])
prompt = get_bounding_box(ground_truth_mask)

# prepare image + box prompt for the model
inputs = processor(image, input_boxes=[[prompt]], return_tensors="pt").to(device)
for k,v in inputs.items():
    print(k,v.shape)
```

```
pixel_values torch.Size([1, 3, 1024, 1024])
original_sizes torch.Size([1, 2])
reshaped_input_sizes torch.Size([1, 2])
input_boxes torch.Size([1, 1, 4])
```

In [42]:

```
sam_model.eval()

# forward pass
with torch.no_grad():
    outputs = sam_model(**inputs, multimask_output=False)
```

In [43]:

```
# apply sigmoid
medsam_seg_prob = torch.sigmoid(outputs.pred_masks.squeeze(1))
# convert soft mask to hard mask
medsam_seg_prob = medsam_seg_prob.cpu().numpy().squeeze()
medsam_seg = (medsam_seg_prob > 0.5).astype(np.uint8)
```

In [44]:

```
def show_mask(mask, ax, random_color=False):
    if random_color:
```

```

        color = np.concatenate([np.random.random(3), np.array([0.6])], axis=0)
    else:
        color = np.array([30/255, 144/255, 255/255, 0.6])
    h, w = mask.shape[-2:]
    mask_image = mask.reshape(h, w, 1) * color.reshape(1, 1, -1)
    ax.imshow(mask_image)


fig, axes = plt.subplots()

axes.imshow(np.array(image))
show_mask(medsam_seg, axes)
axes.title.set_text(f"Predicted mask")
axes.axis("off")

```

Out[44]:

(-0.5, 255.5, 255.5, -0.5)



In [45]:

```


fig, axes = plt.subplots()

axes.imshow(np.array(image))
show_mask(ground_truth_mask, axes)
axes.title.set_text(f"Ground truth mask")
axes.axis("off")

```

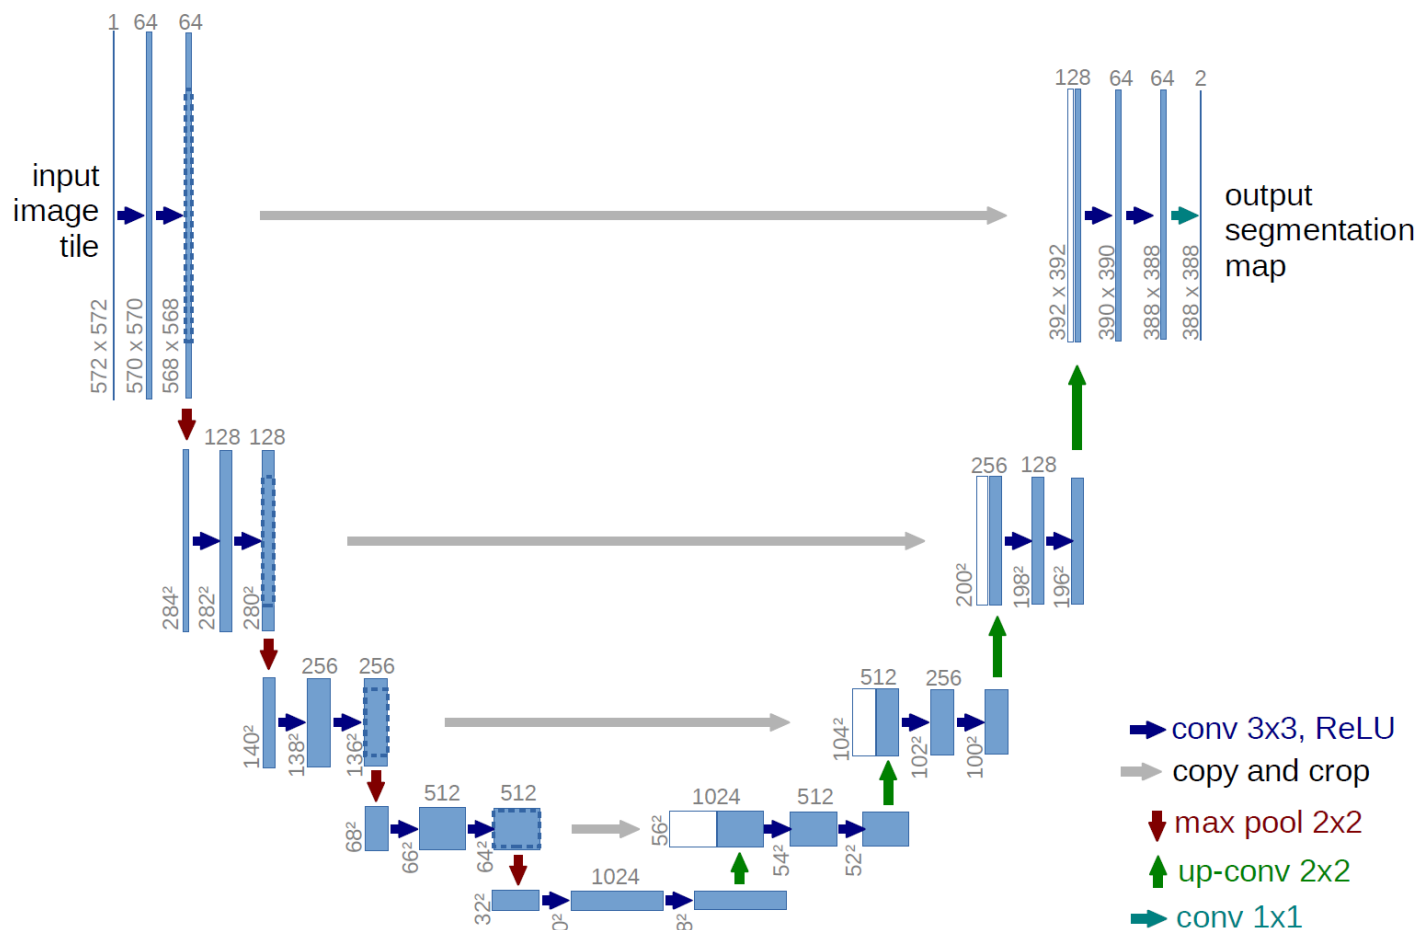
Out[45]:

(-0.5, 255.5, 255.5, -0.5)



RTB AX

UNET modelling



In [47]:

```
class DoubleConv(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(DoubleConv, self).__init__()
        self.conv = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, 3, 1, 1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(inplace=True),
            nn.Conv2d(out_channels, out_channels, 3, 1, 1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(inplace=True),
        )

    def forward(self, x):
        return self.conv(x)
```

In [48]:

```
class UNET(nn.Module):
    def __init__(
        self, in_channels=3, out_channels=1, features=[64, 128, 256, 512],
    ):
        super(UNET, self).__init__()
        self.ups = nn.ModuleList()
        self.downs = nn.ModuleList()
```

```

self.pool = nn.MaxPool2d(kernel_size=2, stride=2)

# Down part of UNET
for feature in features:
    self.downs.append(DoubleConv(in_channels, feature))
    in_channels = feature

# Up part of UNET
for feature in reversed(features):
    self.ups.append(
        nn.ConvTranspose2d(
            feature*2, feature, kernel_size=2, stride=2,
        )
    )
    self.ups.append(DoubleConv(feature*2, feature))

self.bottleneck = DoubleConv(features[-1], features[-1]*2)
self.final_conv = nn.Conv2d(features[0], out_channels, kernel_size=1)

def forward(self, x):
    skip_connections = []

    for down in self.downs:
        x = down(x)
        skip_connections.append(x)
        x = self.pool(x)

    x = self.bottleneck(x)
    skip_connections = skip_connections[::-1]

    for idx in range(0, len(self.ups), 2):
        x = self.ups[idx](x)
        skip_connection = skip_connections[idx//2]

        if x.shape != skip_connection.shape:
            x = TF.resize(x, size=skip_connection.shape[2:])

        concat_skip = torch.cat((skip_connection, x), dim=1)
        x = self.ups[idx+1](concat_skip)

    return self.final_conv(x)

```

In [49]:

```

def test():
    x = torch.randn((3, 1, 161, 161))
    model = UNET(in_channels=1, out_channels=1)
    preds = model(x)
    assert preds.shape == x.shape

```

In [50]:

```

class UNET_Dataset(Dataset):
    def __init__(self, image_dir, mask_dir, transform=None):
        self.image_dir = image_dir
        self.mask_dir = mask_dir
        self.transform = transform
        self.images = os.listdir(image_dir)

    def __len__(self):
        return len(self.images)

    def __getitem__(self, index):
        img_path = os.path.join(self.image_dir, self.images[index])
        mask_path = os.path.join(self.mask_dir, self.images[index] )
        if img_path and mask_path:
            image = np.array(Image.open(img_path).convert("RGB"))
            mask = np.array(Image.open(mask_path).convert("L"), dtype=np.float32)
            mask[mask == 255.0] = 1.0

        if self.transform is not None:

```

```
augmentations = self.transform(image=image, mask=mask)
image = augmentations["image"]
mask = augmentations["mask"]
```

```
return image, mask
```

In [92]:

```
def save_checkpoint(state, filename="my_checkpoint.pth.tar"):
    print("=> Saving checkpoint")
    torch.save(state, filename)

def load_checkpoint(checkpoint, model):
    print("=> Loading checkpoint")
    model.load_state_dict(checkpoint["state_dict"])

def get_loaders(
    train_dir,
    train_maskdir,
    val_dir,
    val_maskdir,
    batch_size,
    train_transform,
    val_transform,
    num_workers=4,
    pin_memory=True,
):
    train_ds = UNET_Dataset(
        image_dir=train_dir,
        mask_dir=train_maskdir,
        transform=train_transform,
    )

    train_loader = DataLoader(
        train_ds,
        batch_size=batch_size,
        num_workers=num_workers,
        pin_memory=pin_memory,
        shuffle=True,
    )

    val_ds = UNET_Dataset(
        image_dir=val_dir,
        mask_dir=val_maskdir,
        transform=val_transform,
    )

    val_loader = DataLoader(
        val_ds,
        batch_size=batch_size,
        num_workers=num_workers,
        pin_memory=pin_memory,
        shuffle=False,
    )

    return train_loader, val_loader

def check_accuracy(loader, model, lst, device="cuda"):
    num_correct = 0
    num_pixels = 0
    dice_score = 0
    model.eval()

    with torch.no_grad():
        for x, y in loader:
            x = x.to(device)
            y = y.to(device).unsqueeze(1)
            preds = torch.sigmoid(model(x))
            preds = (preds > 0.5).float()
            num_correct += (preds == y).sum()
```

```

        num_pixels += torch.numel(preds)
        dice_score += (2 * (preds * y).sum()) / (
            (preds + y).sum() + 1e-8
        )

    print(
        f"Got {num_correct}/{num_pixels} with acc {num_correct/num_pixels*100:.2f}"
    )
    print(f"Dice score: {dice_score/len(loader)}")
    lst.append([dice_score/len(loader), num_correct/num_pixels*100])
    model.train()

def save_predictions_as_imgs(
    loader, model, folder="saved_images/", device="cuda"
):
    model.eval()
    for idx, (x, y) in enumerate(loader):
        x = x.to(device=device)
        with torch.no_grad():
            preds = torch.sigmoid(model(x))
            preds = (preds > 0.5).float()
            torchvision.utils.save_image(
                preds, f"{folder}/pred_{idx}.png"
            )
            torchvision.utils.save_image(y.unsqueeze(1), f"{folder}{idx}.png")

    model.train()

```

In [52]:

```

[!]mkdir saved_images

```

In [103]:

```

# Hyperparameters etc.
LEARNING_RATE = 1e-4
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"
BATCH_SIZE = 20
NUM_EPOCHS = 50
NUM_WORKERS = 0
IMAGE_HEIGHT = 160 # 1280 originally
IMAGE_WIDTH = 240 # 1918 originally
PIN_MEMORY = True
LOAD_MODEL = False
TRAIN_IMG_DIR = "img"
TRAIN_MASK_DIR = "mask"
VAL_IMG_DIR = "img"
VAL_MASK_DIR = "mask"

def train_fn(loader, model, optimizer, loss_fn, scaler):
    loop = tqdm(loader)
    loss = 0

    for batch_idx, (data, targets) in enumerate(loop):
        data = data.to(device=DEVICE)
        targets = targets.float().unsqueeze(1).to(device=DEVICE)

        # forward
        with torch.cuda.amp.autocast():
            predictions = model(data)
            loss = loss_fn(predictions, targets)
            loss += loss

        # backward
        optimizer.zero_grad()
        scaler.scale(loss).backward()
        scaler.step(optimizer)
        scaler.update()

        # update tqdm loop
        loop.set_postfix(loss=loss.item())

```



```
return losses
```

In [104]:

```
UNET_loss = []
train_transform = A.Compose(
    [
        A.Resize(height=IMAGE_HEIGHT, width=IMAGE_WIDTH),
        A.Rotate(limit=35, p=1.0),
        A.HorizontalFlip(p=0.5),
        A.VerticalFlip(p=0.1),
        A.Normalize(
            mean=[0.0, 0.0, 0.0],
            std=[1.0, 1.0, 1.0],
            max_pixel_value=255.0,
        ),
        ToTensorV2(),
    ],)

val_transforms = A.Compose(
    [
        A.Resize(height=IMAGE_HEIGHT, width=IMAGE_WIDTH),
        A.Normalize(
            mean=[0.0, 0.0, 0.0],
            std=[1.0, 1.0, 1.0],
            max_pixel_value=255.0,
        ),
        ToTensorV2(),
    ],)

model_UNET = UNET(in_channels=3, out_channels=1).to(DEVICE)
loss_fn = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model_UNET.parameters(), lr=LEARNING_RATE)

train_loader, val_loader = get_loaders(
    TRAIN_IMG_DIR,
    TRAIN_MASK_DIR,
    VAL_IMG_DIR,
    VAL_MASK_DIR,
    BATCH_SIZE,
    train_transform,
    val_transforms,
    NUM_WORKERS,
    PIN_MEMORY,
)

data = []

if LOAD_MODEL:
    load_checkpoint(torch.load("my_checkpoint.pth.tar"), model_UNET)

print(train_loader, val_loader)
check_accuracy(val_loader, model_UNET, UNET_loss, device=DEVICE)
scaler = torch.cuda.amp.GradScaler()

for epoch in range(NUM_EPOCHS):
    a = train_fn(train_loader, model_UNET, optimizer, loss_fn, scaler)
    data.append(a)

    # save model
    checkpoint = {
        "state_dict": model_UNET.state_dict(),
        "optimizer": optimizer.state_dict(),
    }
    save_checkpoint(checkpoint)

    # check accuracy
    check_accuracy(val_loader, model_UNET, UNET_loss, device=DEVICE)
```

```
# print some examples to a folder
save_predictions_as_imgs(
    val_loader, model_UNET, folder="saved_images/", device=DEVICE
)
```

```
<torch.utils.data.dataloader.DataLoader object at 0x7f96a0b7cd30> <torch.utils.data.dataloader.DataLoader object at 0x7f96a0b7f6d0>
Got 4307618/4992000 with acc 86.29
Dice score: 0.0
```

```
100%|██████████| 7/7 [00:02<00:00, 2.84it/s, loss=0.6]
```

```
=> Saving checkpoint
Got 4307618/4992000 with acc 86.29
Dice score: 0.0
```

```
100%|██████████| 7/7 [00:02<00:00, 2.80it/s, loss=0.531]
```

```
=> Saving checkpoint
Got 4307618/4992000 with acc 86.29
Dice score: 0.0
```

```
100%|██████████| 7/7 [00:02<00:00, 2.67it/s, loss=0.464]
```

```
=> Saving checkpoint
Got 4307618/4992000 with acc 86.29
Dice score: 0.0
```

```
100%|██████████| 7/7 [00:02<00:00, 2.77it/s, loss=0.441]
```

```
=> Saving checkpoint
Got 4307618/4992000 with acc 86.29
Dice score: 0.0
```

```
100%|██████████| 7/7 [00:02<00:00, 2.50it/s, loss=0.371]
```

```
=> Saving checkpoint
Got 4373534/4992000 with acc 87.61
Dice score: 0.2180602252483368
```

```
100%|██████████| 7/7 [00:02<00:00, 2.80it/s, loss=0.459]
```

```
=> Saving checkpoint
Got 4510421/4992000 with acc 90.35
Dice score: 0.5493773221969604
```

```
100%|██████████| 7/7 [00:02<00:00, 2.80it/s, loss=0.44]
```

```
=> Saving checkpoint
Got 4521334/4992000 with acc 90.57
Dice score: 0.604794979095459
```

```
100%|██████████| 7/7 [00:02<00:00, 2.75it/s, loss=0.384]
```

```
=> Saving checkpoint
Got 4449588/4992000 with acc 89.13
Dice score: 0.558542788028717
```

```
100%|██████████| 7/7 [00:02<00:00, 2.81it/s, loss=0.329]
```

```
=> Saving checkpoint
Got 4551908/4992000 with acc 91.18
Dice score: 0.6693931818008423
```

```
100%|██████████| 7/7 [00:02<00:00, 2.52it/s, loss=0.35]
```

```
=> Saving checkpoint
Got 4464802/4992000 with acc 89.44
Dice score: 0.6363086700439453
```

```
100%|██████████| 7/7 [00:02<00:00, 2.81it/s, loss=0.32]
```

```
=> Saving checkpoint
Got 4535576/4992000 with acc 90.86
Dice score: 0.6630691885948181
```

```
100%|██████████| 7/7 [00:02<00:00, 2.80it/s, loss=0.363]
```

=> Saving checkpoint
Got 4560862/4992000 with acc 91.36
Dice score: 0.5571573376655579

100%|██████████| 7/7 [00:02<00:00, 2.83it/s, loss=0.311]

=> Saving checkpoint
Got 4576707/4992000 with acc 91.68
Dice score: 0.585880696773529

100%|██████████| 7/7 [00:02<00:00, 2.83it/s, loss=0.335]

=> Saving checkpoint
Got 4605274/4992000 with acc 92.25
Dice score: 0.6282346844673157

100%|██████████| 7/7 [00:02<00:00, 2.51it/s, loss=0.315]

=> Saving checkpoint
Got 4603880/4992000 with acc 92.23
Dice score: 0.6866751313209534

100%|██████████| 7/7 [00:02<00:00, 2.80it/s, loss=0.373]

=> Saving checkpoint
Got 4606184/4992000 with acc 92.27
Dice score: 0.7114956974983215

100%|██████████| 7/7 [00:02<00:00, 2.69it/s, loss=0.328]

=> Saving checkpoint
Got 4555815/4992000 with acc 91.26
Dice score: 0.6692947149276733

100%|██████████| 7/7 [00:02<00:00, 2.81it/s, loss=0.314]

=> Saving checkpoint
Got 4624328/4992000 with acc 92.63
Dice score: 0.7022212743759155

100%|██████████| 7/7 [00:02<00:00, 2.83it/s, loss=0.271]

=> Saving checkpoint
Got 4637818/4992000 with acc 92.91
Dice score: 0.6904920339584351

100%|██████████| 7/7 [00:02<00:00, 2.61it/s, loss=0.351]

=> Saving checkpoint
Got 4651762/4992000 with acc 93.18
Dice score: 0.7252456545829773

100%|██████████| 7/7 [00:02<00:00, 2.83it/s, loss=0.29]

=> Saving checkpoint
Got 4617021/4992000 with acc 92.49
Dice score: 0.6457382440567017

100%|██████████| 7/7 [00:02<00:00, 2.60it/s, loss=0.313]

=> Saving checkpoint
Got 4585578/4992000 with acc 91.86
Dice score: 0.5727750062942505

100%|██████████| 7/7 [00:02<00:00, 2.84it/s, loss=0.3]

=> Saving checkpoint
Got 4641185/4992000 with acc 92.97
Dice score: 0.6780946850776672

100%|██████████| 7/7 [00:02<00:00, 2.81it/s, loss=0.282]

=> Saving checkpoint
Got 4643256/4992000 with acc 93.01
Dice score: 0.680303156375885

100%|██████████| 7/7 [00:02<00:00, 2.64it/s, loss=0.282]

=> Saving checkpoint
Got 4648877/4992000 with acc 93.13
Dice score: 0.7286368608474731

100%|██████████| 7/7 [00:02<00:00, 2.81it/s, loss=0.297]

=> Saving checkpoint
Got 4692315/4992000 with acc 94.00
Dice score: 0.7388197779655457

100%|██████████| 7/7 [00:02<00:00, 2.54it/s, loss=0.311]

=> Saving checkpoint
Got 4707853/4992000 with acc 94.31
Dice score: 0.7479079365730286

100%|██████████| 7/7 [00:02<00:00, 2.80it/s, loss=0.247]

=> Saving checkpoint
Got 4589345/4992000 with acc 91.93
Dice score: 0.6246888637542725

100%|██████████| 7/7 [00:02<00:00, 2.82it/s, loss=0.246]

=> Saving checkpoint
Got 4699480/4992000 with acc 94.14
Dice score: 0.7438232898712158

100%|██████████| 7/7 [00:02<00:00, 2.68it/s, loss=0.267]

=> Saving checkpoint
Got 4709830/4992000 with acc 94.35
Dice score: 0.7651509046554565

100%|██████████| 7/7 [00:02<00:00, 2.83it/s, loss=0.273]

=> Saving checkpoint
Got 4673356/4992000 with acc 93.62
Dice score: 0.7274051904678345

100%|██████████| 7/7 [00:02<00:00, 2.54it/s, loss=0.281]

=> Saving checkpoint
Got 4682761/4992000 with acc 93.81
Dice score: 0.71808260679245

100%|██████████| 7/7 [00:02<00:00, 2.81it/s, loss=0.246]

=> Saving checkpoint
Got 4682762/4992000 with acc 93.81
Dice score: 0.7282747030258179

100%|██████████| 7/7 [00:02<00:00, 2.82it/s, loss=0.294]

=> Saving checkpoint
Got 4672720/4992000 with acc 93.60
Dice score: 0.7156001925468445

100%|██████████| 7/7 [00:02<00:00, 2.74it/s, loss=0.261]

=> Saving checkpoint
Got 4688685/4992000 with acc 93.92
Dice score: 0.7389267683029175

100%|██████████| 7/7 [00:02<00:00, 2.82it/s, loss=0.287]

=> Saving checkpoint
Got 4691836/4992000 with acc 93.99
Dice score: 0.744437038898468

100%|██████████| 7/7 [00:02<00:00, 2.51it/s, loss=0.342]

=> Saving checkpoint
Got 4689295/4992000 with acc 93.94
Dice score: 0.7644223570823669

100%|██████████| 7/7 [00:02<00:00, 2.80it/s, loss=0.313]

=> Saving checkpoint
Got 4630293/4992000 with acc 92.75
Dice score: 0.7240407466888428

100%|██████████| 7/7 [00:02<00:00, 2.76it/s, loss=0.353]

=> Saving checkpoint
Got 4543496/4992000 with acc 91.02
Dice score: 0.7188430428504944

100%|██████████| 7/7 [00:02<00:00, 2.79it/s, loss=0.301]

=> Saving checkpoint
Got 4595156/4992000 with acc 92.05
Dice score: 0.6678043603897095

100%|██████████| 7/7 [00:02<00:00, 2.82it/s, loss=0.293]

=> Saving checkpoint
Got 4675965/4992000 with acc 93.67
Dice score: 0.7460671067237854

100%|██████████| 7/7 [00:02<00:00, 2.52it/s, loss=0.236]

=> Saving checkpoint
Got 4693329/4992000 with acc 94.02
Dice score: 0.7348936200141907

100%|██████████| 7/7 [00:02<00:00, 2.82it/s, loss=0.246]

=> Saving checkpoint
Got 4697053/4992000 with acc 94.09
Dice score: 0.7307611107826233

100%|██████████| 7/7 [00:02<00:00, 2.72it/s, loss=0.262]

=> Saving checkpoint
Got 4700224/4992000 with acc 94.16
Dice score: 0.7378231883049011

100%|██████████| 7/7 [00:02<00:00, 2.82it/s, loss=0.222]

=> Saving checkpoint
Got 4688433/4992000 with acc 93.92
Dice score: 0.7329518795013428

100%|██████████| 7/7 [00:02<00:00, 2.82it/s, loss=0.22]

=> Saving checkpoint
Got 4718803/4992000 with acc 94.53
Dice score: 0.7581741213798523

100%|██████████| 7/7 [00:02<00:00, 2.55it/s, loss=0.236]

=> Saving checkpoint
Got 4695544/4992000 with acc 94.06
Dice score: 0.7417799234390259

100%|██████████| 7/7 [00:02<00:00, 2.78it/s, loss=0.226]

=> Saving checkpoint
Got 4704884/4992000 with acc 94.25
Dice score: 0.7891114950180054

100%|██████████| 7/7 [00:02<00:00, 2.67it/s, loss=0.25]

=> Saving checkpoint
Got 4755193/4992000 with acc 95.26
Dice score: 0.806406557559967

100%|██████████| 7/7 [00:02<00:00, 2.83it/s, loss=0.221]

=> Saving checkpoint
Got 4721705/4992000 with acc 94.59

Dice score: 0.759075403213501

In [106]:

```
data[0].item()
```

Out[106]:


0.6004032492637634

In [107]:

```
mean_losses = [x.item() for x in data]

plt.plot(list(range(len(mean_losses))), mean_losses)
plt.title('Mean epoch loss')
plt.xlabel('Epoch Number')
plt.ylabel('Loss')

plt.show()
```



In [121]:


```
dice = []
acc = []
for x in unet_loss:
    dice.append(x[0].item())
    acc.append(x[1].item())
```

In [122]:

```
plt.plot(list(range(len(dice))), dice)
plt.title('Mean dice loss')
plt.xlabel('Epoch Number')
plt.ylabel('dice loss')

plt.show()
```






In [123]:

```
plt.plot(list(range(len(acc))), acc)
plt.title('Mean accuracy')
plt.xlabel('Epoch Number')
plt.ylabel('accuracy')

plt.show()
```



In [124]:

```
def plot_images(images, mask = False):
    fig = plt.figure(figsize=(10, 8))
    columns = len(images)
    rows = 1
    for i in range(columns*rows):
        img = images[i]
        fig.add_subplot(rows, columns, i+1)
        if mask:
            img = img > 0.5
        plt.imshow(img, cmap="gray")
```


```
plt.show()
```

In [145]:

```
ground_t = cv2.imread('saved_images/0.png')
plt.imshow(np.array(groud_t))
```

Out[145]:

<matplotlib.image.AxesImage at 0x7f96ac170b20>




In [146]:

```
pred_t = cv2.imread('saved_images/pred_0.png')
plt.imshow(np.array(pred_t))
```

Out[146]:

<matplotlib.image.AxesImage at 0x7f96a089d1e0>



In [156]:

```
for x, y in val_loader:
    x = x.to(device)
```

In [171]:

```
def show_img(loader, model, device="cuda"):
    model.eval()
    for idx, (x, y) in enumerate(loader):
        x = x.to(device=device)
        with torch.no_grad():
            preds = torch.sigmoid(model(x))
            preds = (preds > 0.5).float()

        # Convert tensors to numpy arrays
        y = y.unsqueeze(1)
        preds_np = preds.detach().cpu().numpy()
        y_np = y.detach().cpu().numpy()
        x_np = x.detach().cpu().numpy()

        # Display the predicted and target images
        plt.figure(figsize=(10, 5))
        plt.subplot(1, 3, 1)
        plt.imshow(preds_np[0, 0], cmap='gray')
        plt.title("Predicted")
        plt.axis('off')

        plt.subplot(1, 3, 2)
```



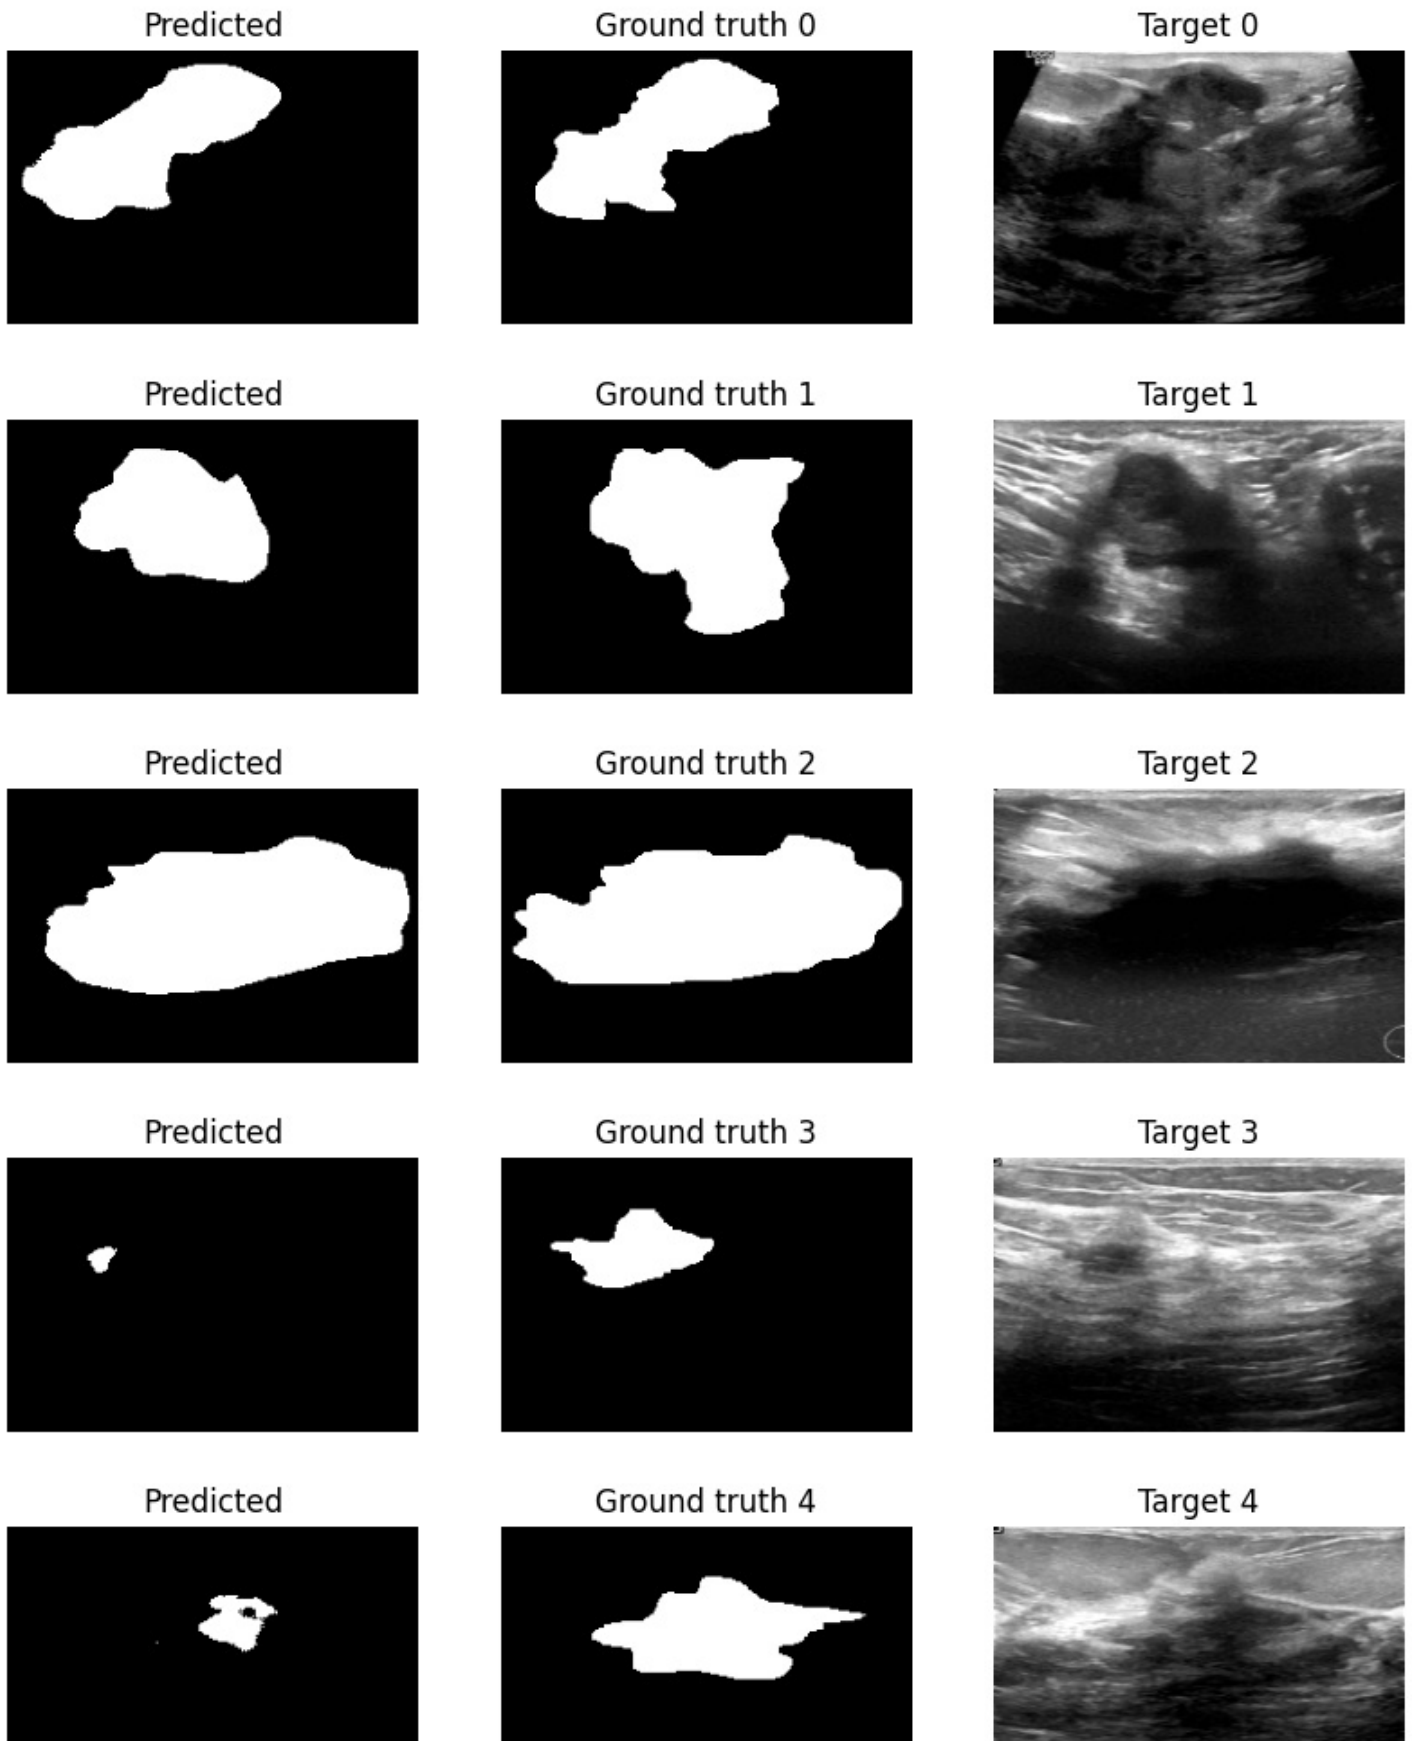
```
plt.imshow(y_np[0, 0], cmap='gray')
plt.title(f"Ground truth {idx}")
plt.axis('off')

plt.subplot(1, 3, 3)
plt.imshow(x_np[0, 0], cmap='gray')
plt.title(f"Target {idx}")
plt.axis('off')

plt.show()
```

In [172]:

```
show_img(val_loader, model_UNET)
```






Predicted



Ground truth 5



Target 5



Predicted



Ground truth 6



Target 6

